

OS Fundamentals & Structure

PART 1 — OS FUNDAMENTALS (refined from your notes)


1. What is an Operating System?

Formal definition

An operating system is system software that manages a computer's hardware resources (CPU, memory, storage, I/O devices) and provides controlled, uniform abstractions and services through which application programs request and use those resources.

Intuition

Hardware is messy and inconsistent — different disks, different network cards, different CPUs, all with their own quirks. Application developers don't want to deal with any of that. The OS's job is to sit in the middle and turn "ugly, device-specific hardware interfaces" into "simple, uniform system calls" that any application can rely on, regardless of what hardware it's actually running on.



It also acts as:

- **Resource manager** — allocates CPU, memory, I/O devices, and files among competing processes.
- **Control program** — supervises execution of user programs to prevent errors and misuse.

Interview answer

"The OS is system software that manages hardware resources and exposes them to applications through a controlled set of abstractions — processes instead of raw CPU time, virtual memory instead of raw RAM, files instead of raw disk blocks — while also enforcing protection between programs."

OS vs Kernel — are they the same?

No — this is one of the most commonly blurred distinctions, and interviewers like to probe it.

- The **kernel** is the core piece of software that runs in privileged (kernel) mode with direct hardware access. It typically handles process scheduling, memory management, and the lowest-level device/interrupt handling.
- The **OS** is the *entire package*: the kernel **plus** everything else needed to make the system usable — shells, system libraries (e.g., glibc), system utilities, system daemons/services, sometimes a GUI, package managers, etc. Many of these run in **user space**, not inside the kernel.

Why do people use the terms interchangeably? Because the kernel is the most architecturally interesting and foundational part — colloquially "the OS" often just means "the kernel + the bare minimum to run it." But technically, Ubuntu and Fedora are different *operating systems* built around the *same* Linux *kernel*.

Mental model:



- **Application layer** — the programs the user actually runs.
- **System calls / OS interfaces** — the boundary applications use to request services; also where user-space OS components (shell, libc, system daemons) live.
- **Kernel** — the privileged core that actually manages hardware.
- **Hardware** — CPU, memory, disks, network cards.

Interview soundbite: "The kernel is the privileged core of the OS that directly manages hardware; the OS is the full software environment — kernel plus user-space libraries, services, and utilities — built around it."

2. Dual Mode Operation

(Kept concise — this is foundational material you already know.)

The CPU runs in exactly one of two modes at any instant, tracked by a **mode bit** in the PSW (Program Status Word):

Mode bit	Mode	Access
1	User mode	Only resources allocated to that process
0	Kernel mode	Full access to any hardware resource

- **User space** — address range allocated to user processes.
- **Kernel space** — address range reserved for the kernel.

Privileged vs non-privileged instructions

- **Privileged instruction** — affects other processes or hardware directly (e.g., changing the mode bit, direct I/O, halting the CPU); executable only in kernel mode.
- **Non-privileged instruction** — affects only the executing process; runs in either mode.

System calls — a controlled doorway, not the only doorway

A system call is how a *program itself* deliberately requests privileged service:

Application → read() → system call → kernel → device/file → result returned

```

User process executing (mode bit = 1)
  |
  | calls system call
  |
  ▼
set mode bit = 0 → switch to kernel mode
  |
  | Execute system call (kernel mode, mode bit = 0)
  |
  | set mode bit = 1 → switch back to user mode
  |
  ▼
Return to user mode (mode bit = 1)

```

Correction to a common misconception: System calls are **not** the only way execution transfers into kernel mode. Two other mechanisms exist:

- **Hardware interrupts** — a device (disk, keyboard, timer) signals the CPU asynchronously, independent of what the running program is doing. The CPU immediately traps to a kernel interrupt handler, regardless of whether the program asked for it.
- **Exceptions (traps)** — synchronous, unplanned events caused by the currently executing instruction itself, such as a divide-by-zero, an invalid memory access (page fault), or an illegal instruction. The CPU traps into the kernel to handle the condition.

System calls are actually implemented *using* the trap mechanism (a deliberate, software-triggered exception) — so all three (syscalls, hardware interrupts, exceptions) fundamentally rely on the same trap/interrupt hardware, but differ in *who* initiates the transfer and *why*.

Why dual mode matters

- **Protection** — a process cannot directly overwrite another process's memory or another user's files.
 - **Isolation** — a crash or bug in one user-space program doesn't directly corrupt kernel state.
 - **Security** — privileged operations (device access, changing page tables) are gated behind kernel-checked entry points.
 - **System stability** — the kernel can enforce invariants (e.g., valid memory ranges, valid file descriptors) before allowing an operation to proceed.
-

3. OS Functionalities

(Kept at existing depth — this is the map of "what an OS does," which becomes essential context for Part 2.)

- Process management
 - CPU scheduling
 - Memory management
 - File-system management
 - I/O management / device management
 - Networking & communication
 - Security / protection / error handling / recovery
 - IPC (Inter-Process Communication)
 - Resource allocation
 - Accounting (tracking resource usage per user/process)
-

4. Evolution of Operating Systems

(Kept concise.)

Stage	Idea
No OS	Manual operation
Uniprogramming (punch cards)	One program at a time
Batch processing	Group similar jobs, run as one batch
Multiprogramming	Multiple programs resident in memory; CPU switches to another when one waits (e.g., on I/O)
Multiprocessing	Multiple CPUs/cores execute tasks genuinely in parallel
Time-sharing	CPU time sliced across processes so multiple interactive users/processes get responsive execution (e.g., Round Robin)
Real-time (RTOS)	Scheduling driven by deadlines, not fairness
Distributed OS	OS layers spread across multiple machines
Modern (mobile/cloud)	Android, cloud-boot systems, etc.

Key distinction interviewers probe:

- **Multiprogramming** — 1 CPU, several programs in memory, CPU *switches* between them. Interleaving, not true simultaneity.
- **Multiprocessing** — >1 CPU, genuinely parallel execution.
- **Time-sharing** — a multiprogramming system tuned for fast turnaround via small time slices, giving interactive users the *illusion* of simultaneity.

5. User Interface Types

- **GUI** — Graphical User Interface (e.g., Windows Explorer).
- **CLI** — Command Line Interface (e.g., bash).
- **Batch** — jobs submitted as a queue, no live interaction (e.g., old-style job scheduling systems, or modern nightly batch jobs).

6. Client-Server Model

- **Client** — initiates a request for a service (e.g., a browser).
- **Server** — the program (running on the same or a different machine) that fulfills the request and sends back a response.
- **Communication** — happens via a defined protocol/interface, not by directly sharing internal state.

Example: a browser (client) sends an HTTP request; a web server (server) processes it and returns a response.

This model becomes directly relevant later: in a **microkernel**, services like file systems and device drivers run as separate user-space *servers*, and applications become *clients* that talk to them via IPC instead of direct function calls.

7. OS Modules

- Process management
- Memory management
- File system
- I/O
- Networking
- Security
- IPC

These are the *conceptual responsibilities* an OS must fulfill. They say nothing about *where* this code physically runs or *how* it's organized — which is exactly the question Part 2 answers.

If an OS has all these responsibilities, the next question is: **how should these components actually be organized?**

PART 2 — OS STRUCTURE AND DESIGN

(New section — taught from first principles, in depth.)

8. Why Does an OS Need a Structure?

An OS is one of the largest, longest-lived pieces of software that exists. It contains:

- process management, CPU scheduling
- memory management, virtual memory
- file systems
- device drivers
- I/O management
- networking
- IPC

- security
- system-call handling

If all of this were dumped together with no organizing principle, the result would be:

- **difficult to understand** — no one could reason about how a change in one part affects another.
- **difficult to maintain** — bug fixes and features risk breaking unrelated functionality.
- **difficult to debug** — a crash could originate almost anywhere.
- **difficult to secure** — no clear boundary between trusted and untrusted code.
- **difficult to modify/extend** — adding a new file system or driver would risk destabilizing the whole system.
- **vulnerable to failures spreading** — one broken component can take down everything sharing its space.

This is why every real OS makes a deliberate architectural choice.

Definition: OS structure (or architecture) describes **how the major components of an operating system are organized, where they execute, and how they communicate with each other.**

The core design questions every OS architecture answers

1. What functionality should be **inside** the kernel (privileged, trusted)?
2. What functionality should be **outside** the kernel (unprivileged, less trusted)?
3. How should components **communicate** with each other?
4. How much **isolation** should components have from one another?
5. How should we balance **performance** against **reliability, security, and maintainability**?

Every architecture you're about to learn — monolithic, microkernel, hybrid, layered, modular — is really just a different *answer* to these five questions.

9. Monolithic Kernel

9.1 What does "monolithic" mean here?


"Monolithic" describes *where code runs and at what privilege level* — **not** how the source code is physically organized.

Clarify what it does **not** mean:

- ❌ It does not mean the kernel is one giant source file.
- ❌ It does not mean there's no internal modularity or organization.
- ❌ It does not mean there are no clean internal interfaces between subsystems.

What it **does** mean: the major OS services — process management, memory management, file systems, networking, device drivers, IPC — all execute together **in the same privileged kernel space**, as one large trusted program, even if that program is internally split into many well-organized source files and subsystems.

9.2 What's inside a monolithic kernel?



Because all of these subsystems live in the same address space at the same privilege level, they can call each other's functions **directly**, the same way any two functions in a normal program call each other — no message passing, no context switch, no crossing a protection boundary.

9.3 Advantages

- **High performance** — a file system requesting memory from the memory manager, or a driver notifying the scheduler, is just a direct function call.
- **Low communication overhead** — no IPC, no context switches between kernel subsystems.
- **Efficient interaction between services** — subsystems can share data structures directly when needed.

9.4 Disadvantages

This is **not** simply "monolithic = unsafe." The real trade-off is about *blast radius*:

- **Large trusted computing base (TCB)** — nearly all OS code (including every built-in driver) runs with full privilege, so all of it must be trusted to be correct.
- **Weaker fault isolation** — because everything shares one address space, a bug in one subsystem (say, a driver) *can* corrupt kernel memory used by an unrelated subsystem (say, the scheduler), since there's no protection boundary between them.
- **A serious kernel bug can potentially bring down the entire system** — not because monolithic kernels are inherently fragile, but because there's no wall stopping a bug in one part from reaching another.
- **Maintaining a large, tightly-coupled kernel is inherently complex** as it grows.

9.5 Example: Linux

Linux is generally described as a **monolithic kernel** — process management, memory management, file systems, and networking all run in kernel space.

Important: Linux is *also* **modular** (see §13). "Monolithic" and "modular" are not opposites — Linux proves that a kernel can run its major services in privileged kernel space (monolithic) *while also* supporting dynamically loadable/unloadable components (modular).

Interview soundbite: "A monolithic kernel places most core OS services in the privileged kernel space, allowing efficient direct communication but creating a larger privileged codebase with weaker internal fault isolation."

10. Microkernel

This is the architecture worth understanding from first principles, since it flips the monolithic design on its head.

The fundamental design question

What if we kept the kernel as small as possible, and moved as many OS services as possible *outside* it, into user space?

Why would anyone want this? Because in a monolithic kernel, *every* piece of kernel code — including third-party drivers — is fully trusted and fully privileged. A bug in one obscure driver can corrupt the memory of the scheduler, the file system, anything. The microkernel philosophy asks: what's the *minimum* amount of code that genuinely *needs* to run with full privilege? Everything else should be pulled out, so that a bug in it can only damage itself, not the whole system.

10.1 What remains inside a microkernel?

Only the bare mechanisms that *truly require* privileged execution — not full policies, just mechanisms. Depending on the specific design, this typically includes:

- Low-level **address-space management** (setting up page tables, the mechanism, not deciding what goes where)
- **Scheduling / thread management** (the mechanism for switching between threads)
- **IPC** (the mechanism processes use to talk to each other)
- Basic **hardware/interrupt handling** (routing interrupts to the right handler)


That's it. Everything else is deliberately pushed out.

10.2 What moves outside the kernel?

Services like:

- File systems
- Networking stacks
- Device drivers
- Other OS services (e.g., authentication)

...run as separate **user-space processes/servers** — ordinary processes, just like any application, except their "job" is to provide an OS service.



Walking through this: an application wants to read a file. It can't call the file system's code directly anymore — the file system isn't part of the kernel, and it isn't part of the application either. It's a *separate process*. So the application sends a **message** (via the microkernel's IPC mechanism) to the file system server, asking it to perform the read. The file system server, in turn, may need to talk to a disk driver — again, a separate process — so *it* also uses IPC. The microkernel's only job in this whole flow is to safely deliver those messages and switch between the processes involved.

10.3 Why IPC becomes central — the causal chain

This is the key insight to internalize, not just memorize:

Services are moved outside the kernel → they become **separate, isolated processes** → separate processes cannot call each other's internal functions directly (they don't share an address space, and there's no privilege escalation without going through the kernel) → therefore they need a **communication mechanism** provided by the kernel → that mechanism is **IPC (message passing)**.

In a monolithic kernel, "the file system needs memory" is a function call. In a microkernel, "the file system needs memory" is a *message* sent to whatever process/mechanism manages memory, which must be delivered, received, processed, and replied to — each of those steps potentially involving a **context switch** and a **copy of data** across the process/kernel boundary.

Analogy: In a monolithic kernel, all departments of a company sit in the same open-plan office and can just shout across the room to each other. In a microkernel, each department is a separate company in a separate building — to get anything done, they have to send formal letters (messages)

through a courier (the microkernel), even for the most trivial questions. It's safer (one company's fire doesn't burn down the others) but slower.

10.4 Advantages

- **Fault isolation** — if the user-space file system server crashes, the kernel itself doesn't crash. In principle, the OS can even restart the failed service.
- **Security** — much less code runs with full privilege; most services run as ordinary, sandboxable user-space processes.
- **Smaller trusted computing base (TCB)** — only the tiny microkernel needs to be trusted for a system-wide guarantee; individual servers can misbehave without compromising everything.
- **Maintainability** — services can be developed, updated, and even replaced independently, without touching the kernel.

10.5 Disadvantages

- **IPC overhead** — every cross-boundary interaction that used to be a function call is now a message send/receive.
- **Extra context switches** — a single high-level operation (e.g., "read a file") may require multiple hops between processes and the kernel.
- **Design complexity** — coordinating many independent servers, defining clean IPC protocols between them, and handling partial failures is genuinely hard to get right.

Important nuance: Do not conclude "microkernels are always slower." This was true of early, naive implementations, but modern microkernel designs (e.g., seL4) have shown that with careful engineering — fast IPC primitives, minimizing unnecessary copies/switches — the overhead can be reduced dramatically. The real trade-off is *engineering effort spent optimizing IPC* vs. the *inherent simplicity* of direct function calls in a monolithic design.

10.6 Examples

- **MINIX** — a teaching/research microkernel OS.
- **QNX** — a commercial, real-time microkernel OS widely used in embedded and safety-critical systems (e.g., automotive).

Interview soundbite: "A microkernel keeps only the minimal privileged mechanisms — IPC, basic scheduling, basic address-space management — inside the kernel, and moves everything else (file systems, drivers, networking) into isolated user-space servers that communicate via message passing. This improves fault isolation and security at the cost of IPC overhead."

11. Hybrid Kernel

Real-world operating systems rarely fit a "pure" category — they're engineering artifacts, optimized for practical constraints, not academic purity.

Basic idea: A hybrid design combines ideas from both monolithic and microkernel approaches — typically keeping some performance-critical services (like the scheduler, memory manager, or even some drivers) running in privileged kernel space for speed, while still using stronger modularity/separation (conceptually, if not always full user-space isolation) for other components.

Why hybrid designs exist: Pure microkernels proved that fault isolation and security are valuable, but paid a real performance cost from IPC overhead in early implementations. OS designers wanted *some* of that isolation benefit without paying the full performance price for genuinely hot-path services — so they picked and chose which pieces to keep privileged.

Trade-off being balanced: performance (favor monolithic) vs. modularity/isolation (favor microkernel) — resolved case-by-case per subsystem instead of applying one rule to the whole kernel.

Examples:

- **Windows NT family** — often described as hybrid; it has a kernel that includes core scheduling/memory management, alongside a design that historically moved some services (like the window manager, at different points) in and out of kernel space for performance reasons.
- **macOS/XNU** — combines the Mach microkernel (message-passing, memory management, scheduling primitives) with BSD kernel components (process model, file system, networking stack) running together in the same kernel space for efficiency.

Note: whether a specific OS "truly" counts as hybrid is sometimes debated by purists — that debate isn't the point for interview purposes. What matters is understanding that **hybrid = a deliberate, pragmatic mix of monolithic-style and microkernel-style choices, made per-subsystem.**

Interview soundbite: "A hybrid kernel pragmatically mixes monolithic and microkernel ideas — keeping performance-critical services in kernel space while applying more modular/isolated design elsewhere. Windows NT and macOS/XNU are commonly cited examples."

12. Layered OS Architecture

This is a **different axis of classification entirely** — it's not "another kind of kernel," it's an *organizational principle* about how to structure software into layers, which can, in principle, be applied within any kernel type.

The idea

Divide the OS into a stack of layers, where each layer:

- provides a well-defined abstraction/interface to the layer above it
- is built using only the services of the layer(s) below it

```
Layer 5 → User Applications
Layer 4 → User Services
Layer 3 → File / I/O Services
Layer 2 → Process / Memory Management
Layer 1 → Hardware Abstraction
Layer 0 → Hardware
```

- **Higher layers depend on lower layers** — e.g., the file/I/O layer depends on process/memory management being available underneath it.
- **Lower layers should not depend on higher layers** — the hardware abstraction layer shouldn't need to know anything about file services above it. This one-directional dependency is what makes the system reason-able: you can understand layer 2 completely without knowing anything about layer 3, 4, or 5.

Key benefit: each layer provides an abstraction to the layer above it, hiding its own internal complexity — the same principle as the OS-as-a-whole hiding hardware complexity from applications, just applied recursively *inside* the OS.

Advantages

- **Modularity** — each layer is a self-contained unit of understanding.
- **Easier reasoning/debugging** — a bug can often be localized to "somewhere in layer N" based on which abstraction is misbehaving.
- **Clearer interfaces** — each layer boundary is a well-defined contract.
- **Easier testing** — layers can be tested against mocked/stubbed lower layers.
- **Controlled dependencies** — the one-directional dependency rule prevents circular, tangled coupling.

Disadvantages

- **Choosing correct layer boundaries is hard** — draw them wrong, and you end up with awkward abstractions or layers that barely hide anything.
- **Strict layering can introduce overhead** — if every request must pass through every layer sequentially, that's more work than a direct shortcut would be.
- **Some operations naturally cross multiple layers** — e.g., a page fault touches memory management *and* the file system (if paging to disk) *and* I/O, and strict layering can make this awkward to express cleanly.
- **Real OSes rarely follow perfectly strict layering** in practice — pragmatic shortcuts are common.

Layered architecture is primarily a **design/organizational concept**, not another item in the "monolithic vs. microkernel" list. A monolithic kernel's internals can be organized in layers; a microkernel's servers can each be internally layered too.

13. Modular Kernel

Easy to confuse with "monolithic" — but it answers a *different* question.

Monolithic answers: *where does the code run* (privileged kernel space vs. unprivileged user space)?

Modular answers: *how is the kernel's code packaged and loaded* (built permanently into the kernel image vs. dynamically loadable at runtime)?

Definition: A modular kernel can dynamically **load or unload** certain kernel components — most commonly device drivers or file-system drivers — **without requiring the entire kernel to be rebuilt or the system to be rebooted.**

Loadable Kernel Modules (LKMs)

- A device driver, for example, can be compiled as a separate module (`.ko` file on Linux) and loaded into the running kernel only when that specific hardware is present (`insmod` / `modprobe`), then unloaded when no longer needed (`rmmmod`).
- This means the kernel doesn't need to ship with *every possible driver* compiled permanently into it — it can load exactly what the current machine needs, when it needs it.

Why this is useful:

- **Extensibility** — new hardware support can be added without recompiling the whole kernel.
- **Maintainability** — a buggy driver module can be updated and reloaded independently.
- **Reduced memory footprint** — unused modules simply aren't loaded.

Crucial distinction:

	Monolithic	Modular
Answers	<i>Where</i> does code run?	<i>How</i> is code packaged/loaded?
Contrast is with	Microkernel	Kernel built as one static, non-loadable image

A kernel can be both monolithic and modular at the same time — and Linux is the textbook example: its core services (process/memory management, file systems, networking) run in kernel space (monolithic), while drivers and some file-system implementations can be compiled as loadable modules (modular). These are two independent design dimensions, not opposing categories.

14. Compare OS Structures

Aspect	Monolithic	Microkernel	Hybrid	Layered
Main idea	Run all major services in kernel space for speed	Keep kernel minimal; push services to user-space servers for isolation	Pragmatically mix: keep some services privileged, isolate others	Organize software into abstraction layers with one-directional dependencies
Kernel scope	Large (most OS services included)	Very small (IPC, basic scheduling, basic memory/interrupts only)	Medium — larger than a pure microkernel, smaller than everything	Not about kernel size — orthogonal organizational principle
Where services run	Kernel space	User space (as servers), talking via IPC	Mixed — some in kernel space, some isolated	Depends on which layer; not defined by this axis
Communication	Direct function calls	Message passing (IPC)	Mostly direct calls for privileged parts; may use IPC for isolated parts	Calls flow strictly between adjacent layers
Performance	High (no IPC/context-switch overhead between subsystems)	Historically lower due to IPC overhead; modern designs narrow this gap	Aims for near-monolithic performance on hot paths	Can add overhead if every call traverses every layer
Fault isolation	Weak — a bug can corrupt unrelated kernel state	Strong — a failed server doesn't crash the kernel	Mixed — strong for isolated parts, weak for privileged parts	Not directly about fault isolation — about code organization
Security implications	Larger trusted computing base	Smaller TCB; less code is fully privileged	TCB size depends on what's kept privileged	Orthogonal — layering doesn't by itself change privilege boundaries
Complexity	Complex due to tight coupling as it grows	Complex due to IPC protocol design & coordination	Complex due to case-by-case design decisions	Complex to get layer boundaries right
Main advantage	Speed, simplicity of direct calls	Reliability, security, smaller TCB	Balances performance and isolation	Clean reasoning, easier maintenance/testing

Aspect	Monolithic	Microkernel	Hybrid	Layered
Main disadvantage	Weak isolation, large TCB	IPC overhead, coordination complexity	Classification/purity debates; still nontrivial complexity	Rigid boundaries can be awkward or costly
Example	Linux	MINIX, QNX	Windows NT, macOS/XNU	Historically: THE OS (academic); conceptually used within many modern systems


The actual trade-offs, in words

There is no universally "best" architecture — each is optimized for different priorities:

- **Monolithic wins when raw performance and simplicity of implementation matter most**, and you're willing to accept that a bug in a driver could theoretically corrupt kernel state elsewhere. This describes most general-purpose desktop/server OSes (Linux, and historically most Unix systems), where performance at scale matters enormously and the ecosystem has matured enough that catastrophic driver bugs are rare in practice.
- **Microkernel wins when reliability and fault containment matter more than raw throughput** — safety-critical embedded systems (QNX in automotive/medical devices) where a driver crash must never bring down the whole system, even at some performance cost.
- **Hybrid exists because most real systems need to ship a general-purpose OS that's competitive on performance benchmarks, while still wanting *some* of the isolation benefits** microkernels pioneered — so they cherry-pick.
- **Layered is not competing with the other three at all** — it's a *complementary* organizing principle that any of the above can (and often do) apply internally to manage complexity, regardless of where code ultimately executes.

15. Connecting OS Structure With Dual Mode — the unified mental model

This ties everything together. Every architectural decision in Part 2 is really just a different answer to: **"How much of the OS runs in kernel mode, versus user mode, and how do the pieces talk to each other?"**



Walking through this with a concrete example — **a microkernel handling a file read:**

1. The application calls `read()` — a **system call**, transferring control from user mode into the **microkernel** (mode bit flips to 0).
2. The microkernel doesn't implement file systems itself — it only knows how to route messages. It delivers an **IPC message** to the file system server (a user-space process).
3. Delivering that message and switching to the file system server process involves a **context switch** — the mode bit flips back to 1 (user mode), but now executing the file system server's code instead of the application's.
4. If the file system server needs to talk to a disk driver (also a separate user-space process), it sends *another* IPC message, causing another switch.
5. Eventually, a **response** flows back the same way, and the application resumes.

Now contrast with a **monolithic kernel handling the same read:**

1. The application calls `read()` — a **system call**, transferring control into the **kernel** (mode bit flips to 0).
2. The kernel's file system code runs *directly*, in the same address space and privilege level as the process/memory management and driver code.
3. It calls the disk driver's function *directly* — no message, no extra context switch.
4. Control returns to user mode once, at the end.

Same system call, same dual-mode mechanism underneath — but the *number of privilege/process boundary crossings* differs dramatically depending on the architecture. This is the single idea that unifies everything in this chapter: **architecture determines how many boundaries a given operation must cross, and dual-mode operation is the hardware mechanism that makes crossing any boundary safe.**

Final Interview Questions With Answers

Q1. What is an Operating System?

Answer: System software that manages hardware resources (CPU, memory, storage, I/O) and exposes them to applications through controlled, uniform abstractions (processes, virtual memory, files) while enforcing protection between programs.

If interviewer asks further: Be ready to name the three lenses — resource manager, control program, and abstraction/service provider.

Q2. What is the difference between an OS and a kernel?

Answer: The kernel is the privileged core that directly manages hardware (scheduling, memory, low-level I/O) and runs in kernel mode. The OS is the complete package — the kernel plus user-space components like shells, system libraries, daemons, and utilities.

Key distinction: "Kernel" describes *a piece of privileged software*; "OS" describes *the entire system built around it*.

Q3. Why do we need dual-mode operation?

Answer: To protect the system: if any user program could execute privileged instructions freely, a bug or malicious program could corrupt other processes' memory, crash the system, or bypass security. Dual mode enforces a hardware-checked boundary so only trusted, controlled code paths (the kernel) can touch sensitive hardware/state.

Q4. Difference between user mode and kernel mode?

Answer: User mode restricts a process to its own allocated resources and disallows privileged instructions; kernel mode allows full hardware access and privileged instruction execution. The CPU's mode bit (in the PSW) tracks which mode is active.

Q5. What is a system call?

Answer: A controlled, software-triggered mechanism (a type of trap) that lets a user-mode program request a specific privileged service from the kernel — e.g., `read()`, `write()`, `fork()` — causing a mode switch into the kernel, execution of the requested service, and a switch back.

If interviewer asks further: Mention that interrupts and exceptions can *also* cause a switch into kernel mode — system calls aren't the only trigger.

Q6. What is a monolithic kernel?

Answer: A kernel design where the major OS services — process management, memory management, file systems, networking, drivers — all execute together in privileged kernel space, allowing fast direct function calls between them but sharing fate in the same trust/privilege domain.

Q7. Monolithic vs microkernel — what's the core trade-off?

Answer: Monolithic keeps services in kernel space for speed (direct calls) at the cost of weaker fault isolation (a bug anywhere can affect the whole kernel). Microkernel keeps only minimal mechanisms (IPC, scheduling, basic memory) privileged and moves everything else into isolated user-space servers, trading some performance (IPC overhead) for much stronger fault isolation and a smaller trusted computing base.

Q8. Why does a microkernel rely so heavily on IPC?

Answer: Because moving services out of the kernel turns them into separate, isolated user-space processes. Isolated processes can't call each other's internal functions directly — there's no shared address space or elevated privilege to do so safely. So the kernel must provide a controlled communication mechanism — IPC (message passing) — for these services to cooperate at all.

Q9. What are the advantages and disadvantages of microkernels?

Answer: Advantages: fault isolation (a crashed server doesn't crash the kernel), a smaller trusted computing base, better security (less code is fully privileged), and independent maintainability of services. Disadvantages: IPC/context-switch overhead for operations that would be a single function call in a monolithic design, and the engineering complexity of coordinating many independent servers.

Key distinction: This overhead is an *implementation cost*, not an inherent law — well-optimized microkernels (e.g., seL4) narrow the gap significantly.

Q10. What is a hybrid kernel?

Answer: A pragmatic design that mixes monolithic and microkernel ideas — keeping some performance-critical services in privileged kernel space while applying more modular or isolated design to others. Windows NT and macOS/XNU are commonly cited examples.

Q11. What is layered OS architecture?

Answer: An organizational principle where the OS is divided into layers, each providing an abstraction to the layer directly above it and depending only on the layer(s) below it — never the reverse. It improves modularity, debuggability, and testability, at the cost of potential overhead and difficulty in choosing correct boundaries.

Key distinction: Layering is an *organizational axis*, orthogonal to whether the OS is monolithic or microkernel-based — it's not a third "kind of kernel."

Q12. Monolithic vs modular kernel — what's the actual difference?

Answer: Monolithic is about *where code executes* (privileged kernel space, as opposed to user-space in a microkernel). Modular is about *how code is packaged and loaded* (as dynamically loadable components, as opposed to being permanently built into a static kernel image). They answer different questions entirely.

Q13. Can Linux be both monolithic and modular? How?

Answer: Yes. Linux is monolithic because its core services (process management, memory management, file systems, networking) execute in kernel space. It is also modular because many components — especially device drivers and some file-system implementations — can be compiled as Loadable Kernel Modules (`.ko` files) and inserted/removed from the running kernel without a reboot or full rebuild. "Monolithic" and "modular" describe independent dimensions, not opposing categories.

Q14. How does OS structure affect security, reliability, and performance?

Answer: Structure determines how much code is fully privileged (the trusted computing base) and how isolated components are from each other. More code in kernel space (monolithic) generally means higher performance (direct calls) but a larger attack surface and weaker fault containment. More isolation (microkernel) means smaller TCB and better fault containment, at the cost of communication overhead. Hybrid and layered approaches are attempts to get favorable trade-offs on a case-by-case or organizational basis rather than applying one blanket rule.

Q15. Scenario: If a device driver crashes, how would the impact differ between a monolithic and a microkernel design?

Answer (model reasoning):

- In a **monolithic** kernel, the driver's code runs in the same address space and privilege level as the rest of the kernel. If it crashes (e.g., dereferences a bad pointer), it can corrupt kernel memory used by *other* subsystems — the scheduler, the file system, anything — and typically brings down the entire system (a full kernel panic/crash), because there's no protection boundary between the driver and the rest of the kernel.
- In a **microkernel**, the driver runs as an isolated user-space process, separate from the kernel and from other servers. If it crashes, the microkernel (and other unrelated servers) keep running. The specific service the driver provided becomes unavailable, but — depending on the system's design — it may even be possible to detect the crash and **restart just that driver process**, without rebooting the machine or affecting unrelated functionality.
- The underlying reason for the difference is exactly the isolation boundary discussed in §10.4/§14: monolithic shares one fault domain across everything; microkernel gives each service its own fault domain, at the cost of needing IPC to talk between them.